

A Bayesian Approach to Merger Review

Matt Panhans & Charles Taragin

2026-02-01

1 Introduction

Merger review inherently takes place in a data-constrained environment. Antitrust agencies and practitioners are frequently tasked with simulating the effects of a proposed combination using only a single cross-section of market shares, prices, and potentially a few margin observations. In this setting, standard econometric techniques that rely on asymptotics are fundamentally inapplicable. There is simply not enough data to support traditional estimation.

This constraint forces valid structural modeling to confront a deeper, often ignored issue in Industrial Organization: **Model Uncertainty**. Theoretical ambiguity is a genuine feature of many markets. For example, in industries with complex institutional features—such as banking or healthcare—it is often unclear whether firms compete purely on price (Bertrand), localized monopoly power (Monopolistic Competition), or through sealed-bid mechanisms (Auctions). While the field has developed valid tests for non-nested model selection (e.g., Rivers and Vuong (2002)), these procedures typically rely on asymptotic properties that may not hold in a merger review setting with limited degrees of freedom. Moreover, such tests generally seek to identifying a single “best” model. In finite samples where the data may be ambiguous, this forces a binary choice that ignores the remaining uncertainty.

To overcome this, practitioners typically resort to **calibration methods**. This approach essentially treats the structural equations of a specific competitive model (e.g., Bertrand) as exact constraints to back out unobserved parameters like marginal costs or demand elasticities. Effectively, calibration is a “poor man’s Bayesian analysis”—it is equivalent to Bayesian inference with degenerate priors that assign 100% probability to a specific theoretical model and 0% to all others.

The fragility of this approach is exposed when the “true” model of competition is unknown. As demonstrated in Panhans and Taragin (2023), conditioning on different standard models yields contradictory predictions about merger effects. By committing to a single model, the analyst implicitly assumes away all structural uncertainty, leading to overconfident policy prescriptions.

In this paper, we propose replacing this rigid approach with the **full Bayesian machinery**. Unlike frequentist tests that require asymptotics to reject a model, the Bayesian framework is designed “from the ground up” to quantify uncertainty through probability. By explicitly modeling the joint distribution of observed data and parameters, we address the “model selection problem” directly through **Bayesian Model Averaging (BMA)**. Instead of forcing a binary choice between statistically indistinguishable theories, we estimate the posterior probability of each model given the finite data available.

Our contribution is not the development of new Bayesian statistical theory, but rather the application of these standard tools to resolve a persistent identification failure in antitrust policy. While BMA is well-established in macroeconomics and finance, its utility for merger simulation has been overlooked. We show that BMA provides the missing link between the theoretical ambiguity of firm conduct and the practical necessity of reaching a decision in data-constrained reviews. By operationalizing this “conduct-agnostic” framework, we offer a rigorous alternative to the ad-hoc assumptions that currently dominate imperfectly competitive market analysis.

2 Structural Models of Competition

We consider four canonical models of competition. We assume a standard Logit demand system where the market share s_j of product j is given by:

$$s_j(\mathbf{p}) = \frac{\exp(\delta_j)}{\sum_{k=0}^J \exp(\delta_k)}$$

where $\delta_j = \beta x_j - \alpha p_j + \xi_j$ is the mean utility (net of price) and α is the price sensitivity parameter. The outside good share is denoted as s_0 .

2.1 Why Share-Consistent Demand?

We deliberately adopt a demand system consistent with share-based concentration metrics for three reasons rooted in the institutional reality of antitrust enforcement.

First, **legal doctrine operates on market shares**. From the *Philadelphia National Bank* structural presumption to the modern *Merger Guidelines*, agencies primarily screen and challenge mergers based on share-based concentration measures (e.g., HHI). A demand model that does not map market structure effectively to price effects would be theoretically flexible but institutionally irrelevant. Among demand systems consistent with this share-centric paradigm, Logit and CES are the primary tractable options.

Second, **data constraints preclude flexibility**. While Random Coefficients Logit allows for richer substitution patterns, it requires micro-data that is rarely available in the initial stages of merger review. In the absence of such data, flexible models collapse back to calibration on shares, effectively mimicking the Logit outcome but with less transparency.

Third, our Bayesian approach **endogenizes market definition**. A common critique of Logit is its sensitivity to the definition of the outside good. Unlike standard calibration—which fixes s_0 ad-hoc—we treat the outside share s_0 as an unknown parameter to be estimated. This allows the model to probabilistically “size the market” based on the observed tension between prices and shares, absorbing the uncertainty of market definition directly into the posterior.

In our Bayesian framework, we estimate a structural parameter $\tilde{\alpha}$ that scales the margin relationships derived below.

2.2 Bertrand Competition

...

3 Bayesian Methodology

3.1 Model Estimation

We estimate each model $m \in \{\text{Cournot}, \text{Bertrand}, \text{Auction}, \text{MonComp}\}$ using a Bayesian hierarchical framework implemented in Stan. For each model, we specify the joint likelihood of observed market shares and margins. The structural parameters—including the price sensitivity α , marginal costs c , and the **outside share** s_0 —are estimated simultaneously. Unlike standard calibration approaches that often “plug in” an arbitrary value for the outside good (e.g., 50%), we treat s_0 as an unknown parameter to be inferred from the data (conditioned on priors), further capturing the uncertainty inherent in market definition.

3.2 Bertrand Competition

Following the notation in Panhans and Taragin (2023), under Bertrand competition with differentiated products, the margin is given by:

$$m_j \propto \frac{1}{1 - s_j(1 - s_0)}$$

This form captures the standard internalization of cross-elasticities in a Logit setting.

3.3 Second-Score Auction

In the Second-Score Auction model (often applicable to bidding markets), the relationship between the margin and market shares is derived as:

$$m_j \propto \frac{-\log(1 - s_j(1 - s_0))}{s_j(1 - s_0)}$$

This model captures competitive dynamics where the pricing is determined by the distribution of the second-best valuation.

3.4 Cournot Competition

For Cournot competition (quantity setting) under Logit demand, the margin relates to the outside good share as:

$$m_j \propto \frac{s_0 + s_j}{s_0} = 1 + \frac{s_j}{s_0}$$

This relationship implies that as a firm's market share increases (relative to the outside option), its margin increases above the competitive baseline.

3.5 Monopolistic Competition

Finally, we include a model of Monopolistic Competition. In this limiting case, firms have market power but do not strategically interact with specific rivals. This implies a constant scaling:

$$m_j \propto 1$$

This serves as a baseline where margins are determined solely by own-price elasticity, independent of local market share variations.

4 Theoretical Intuition

The core intuition behind our identification strategy relies on the fact that different models of competition predict different margins for the same observed market shares and prices. Since we treat margins as observed data in our Bayesian framework, the model that predicts margins closest to the observed values will receive higher posterior probability.

We formalize this intuition with the following proposition, which establishes a rank ordering of the predicted margins across the candidate models.

Proposition 1 (Structural Margin Hierarchy). *Conditioning on market shares \mathbf{s} and a given price sensitivity α , the predicted margins for firm j satisfy the strict ordering:*

$$m_j^{\text{Cournot}} > m_j^{\text{Bertrand}} > m_j^{\text{Auction}} > m_j^{\text{MonComp}}$$

Proof. Using the margin relationships derived above: 1. **Cournot:** $m_j \propto 1 + s_j/s_0$ 2. **Bertrand:** $m_j \propto (1 - s_j(1 - s_0))^{-1} \approx 1 + s_j$ 3. **Auction:** $m_j \propto -\log(1 - s_j(1 - s_0))/(s_j(1 - s_0)) \approx 1 + s_j/2$ 4. **MonComp:** $m_j \propto 1$

For strictly positive market shares and $s_0 < 1$, the Cournot markup factor (scaled by $1/s_0$) is strictly the largest, followed by Bertrand (differentiated scaling), Auction (second-score scaling), and Monopolistic Competition (constant).

Corollary 1 (Rank Ordering of Calibrated Alphas). *To rationalize a fixed set of observed margins, shares, and prices, the calibrated price sensitivity parameters α must satisfy the same strict ordering:*

$$\alpha^{\text{Cournot}} > \alpha^{\text{Bertrand}} > \alpha^{\text{Auction}} > \alpha^{\text{MonComp}}$$

Bayesian Identification Strategy. This corollary illuminates the identification mechanism. Each model “prefers” a different region of the parameter space to explain the supply-side data (margins). However, the demand-side data (prices and shares) provides an independent anchor for α .

In our Bayesian framework, the models compete to explain the joint distribution of data. The posterior probability favors the model that minimizes the tension between the two data sources—effectively, the model that **“does the best job” of predicting the observed margins** given the constraints imposed by the demand system. For example, if the demand data pin down a moderate α , but the accounting margins are high, the Cournot model (which predicts high margins naturally) will have a higher likelihood than the Bertrand or Auction models, which would require an implausibly low α (high markup) to fit the same data.

5 Bayesian Methodology

5.1 Model Estimation

We estimate each model $m \in \{\text{Cournot}, \text{Bertrand}, \text{Auction}, \text{MonComp}\}$ using a Bayesian hierarchical framework implemented in Stan. For each model, we specify the joint likelihood of observed market shares and margins. The structural parameters (including α and costs) are estimated simultaneously.

5.2 Model Comparison and Averaging

We account for model uncertainty through Bayesian Model Averaging (BMA). We compute the posterior model probability $P(M_k|D)$ for each model k :

$$P(M_k|D) = \frac{P(D|M_k)P(M_k)}{\sum_j P(D|M_j)P(M_j)}$$

where $P(D|M_k)$ is the marginal likelihood of the data under model k . We typically assume uniform prior probabilities $P(M_k)$.

Predictions (e.g., merger price effects Δp) are then averaged:

$$E[\Delta p|D] = \sum_k E[\Delta p|D, M_k]P(M_k|D)$$

6 Bayesian Implementation

The models are estimated using the HMC NUTS sampler in Stan. * **Likelihood:** We model the joint distribution of demand (log-shares) and supply (inverse margins). * **Priors:** We use weakly informative priors for structural parameters. * **Computation:** We approximate the marginal likelihoods using bridge sampling or LOO-CV to compute Bayes Factors.

7 Application: The Banking Industry

We apply our Bayesian framework to the U.S. banking industry, a sector that is foundational to modern antitrust enforcement. The “structural presumption”—the legal doctrine that high market concentration is presumptive evidence of anticompetitive harm—originated in the Supreme Court’s seminal *United States v. Philadelphia National Bank* (1963) decision. Despite this long history, the industry remains a subject of intense debate regarding the appropriate model of competition. Are banks setting deposit rates (Bertrand), competing for a quantity of loanable funds (Cournot), or bidding for customer relationships (Auctions)?

Aside from its legal significance, banking offers a unique empirical laboratory for testing structural models. As noted by Panhans and Taragin (2023), it is one of the few environments where researchers can observe the full triad of necessary data—prices (interest rates), quantities (deposit shares), and margins (net interest margins)—continuously over long time horizons.

- **Mergers:** The industry has undergone massive consolidation, providing a rich set of “natural experiments” for merger simulation.
- **Data Availability:** Unlike most industries where marginal costs are unobservable, banking accounting data (Call Reports) allows for the construction of reliable proxies for variable costs and margins.
- **Local Markets:** The clear definition of geographic markets (e.g., Federal Reserve banking markets) allows for cross-sectional variation that acts as a powerful source of identification.

This combination of data richness and theoretical ambiguity makes banking the ideal setting to demonstrate the power of Bayesian Model Averaging.

8 Case Study: United States v. Philadelphia National Bank (1963)

To demonstrate the practical utility of our Bayesian framework, we apply it to the data from the landmark *United States v. Philadelphia National Bank* (1963) case. This case established the “structural presumption” in antitrust law, yet the nature of competition in this market remains a subject of theoretical interest.

8.1 Data Overview

We utilize the bank-level data from the period immediately preceding the merger (1960). The sample consists of 39 banks competing in the local Philadelphia market. Unlike stylized examples, the observed data reveals significant heterogeneity in market shares, deposit rates, and operating margins.

Table 1: Summary Statistics: PNB Case Data (1960)

Variable	Mean	SD	Min	Max
Margin	1.827	0.340	1.142	2.321
shareIn	0.084	0.079	0.011	0.230
rate_deposits	2.290	0.429	1.128	2.733

8.2 Model Comparison Results

We run our Bayesian Model Comparison on this single-market dataset using a **Strict Structural Specification** (mimicking standard antitrust calibration). We remove the supply intercept to force α to match the margin levels, and we use random effects to fit demand shares.

The results below show the comparison. Note that while **Monopolistic Competition** is favored, the overall fit is constrained by the tension between the margin-implied elasticity and the observed price-share correlations.

Table 2: Bayesian Model Comparison: PNB Case

Model	LogML	LOOIC	RMSE	Alpha
MonCom	-1.137035e+66	-52.96	0	0.25
Cournot	-3.608157e+223	-34.76	0	0.34
Bertrand	-2.476595e+268	-38.46	0	0.27
Auction	-1.986719e+291	-36.21	0	0.26

The Bayes Factor favoring Monopolistic Competition is substantial. This suggests that despite the “structural presumption” linking shares to market power, banks in the pre-merger Philadelphia market effectively acted as local monopolists—setting margins based on customer inelasticity rather than engaging in strategic interaction with local rivals.

8.3 Parameter Stability

A key advantage of our framework is that it estimates structural parameters simultaneously with model selection. As shown below, the estimates for demand elasticity (α) and the outside share (s_0) remain stable across supply models, indicating robust identification of the demand system despite the uncertainty in conduct.

Table 3: Parameter Estimates (Mean and SD) vs Priors

Parameter	Prior	Bertrand	Auction	Cournot	MonCom
Alpha	1.13 (0.60)	0.269 (0.020)	0.259 (0.019)	0.338 (0.043)	0.254 (0.016)
S0	0.33 (0.20)	0.108 (0.030)	0.109 (0.030)	0.143 (0.043)	0.120 (0.033)
Sigma_Share	0.00 (1.10)	0.043 (0.029)	0.042 (0.030)	0.038 (0.031)	0.022 (0.030)
Sigma_Margin	0.00 (0.13)	0.140 (0.038)	0.141 (0.038)	0.201 (0.051)	0.138 (0.031)
Rho	0.00 (0.25)	-0.002 (0.449)	-0.020 (0.442)	0.117 (0.459)	0.390 (0.528)
Cutoff	0.03 (0.02)	NA (NA)	NA (NA)	NA (NA)	NA (NA)

9 Conclusion

We have presented a Bayesian framework for merger review that explicitly accounts for model uncertainty. By shifting the focus from selecting a single “best” model to averaging across a set of plausible candidate theories, we offer a more robust path for antitrust enforcement.

9.1 Policy Implications

Our results suggest that the current practice of relying on a single structural model (often Bertrand) may lead to overconfident and potentially biased predictions of merger effects. In the banking sector application, we find that BMA often yields more conservative harm estimates than the pure Bertrand baseline but validates risk in areas where alternative models (like Cournot) might otherwise be overlooked. This approach provides a rigorous, data-driven path to resolving model ambiguity in individual merger cases, moving beyond the “all-or-nothing” nature of current enforcement debates.

9.2 Future Directions

While this paper focuses on case-specific adjudication, a natural extension is to apply this BMA framework to the design of **Merger Screening Thresholds**. Future work could use the galaxy of BMA-weighted predictions to derive “conduct-robust” concentration screens (e.g., HHI levels) that minimize decision error across a portfolio of uncertain conduct types. Additionally, extending the framework to include more flexible demand specifications (e.g., Random Coefficients Logit) would further validate its utility.

10 References

- Panhans, Matthew T., and Charles Taragin. 2023. “Consequences of Model Choice in Predicting Horizontal Merger Effects.” *International Journal of Industrial Organization* 89: 102986. <https://doi.org/10.1016/j.ijindorg.2023.102986>.
- Rivers, Douglas, and Quang Vuong. 2002. “Model Selection Tests for Nonlinear Dynamic Models.” *The Econometrics Journal* 5 (1): 1–39.